

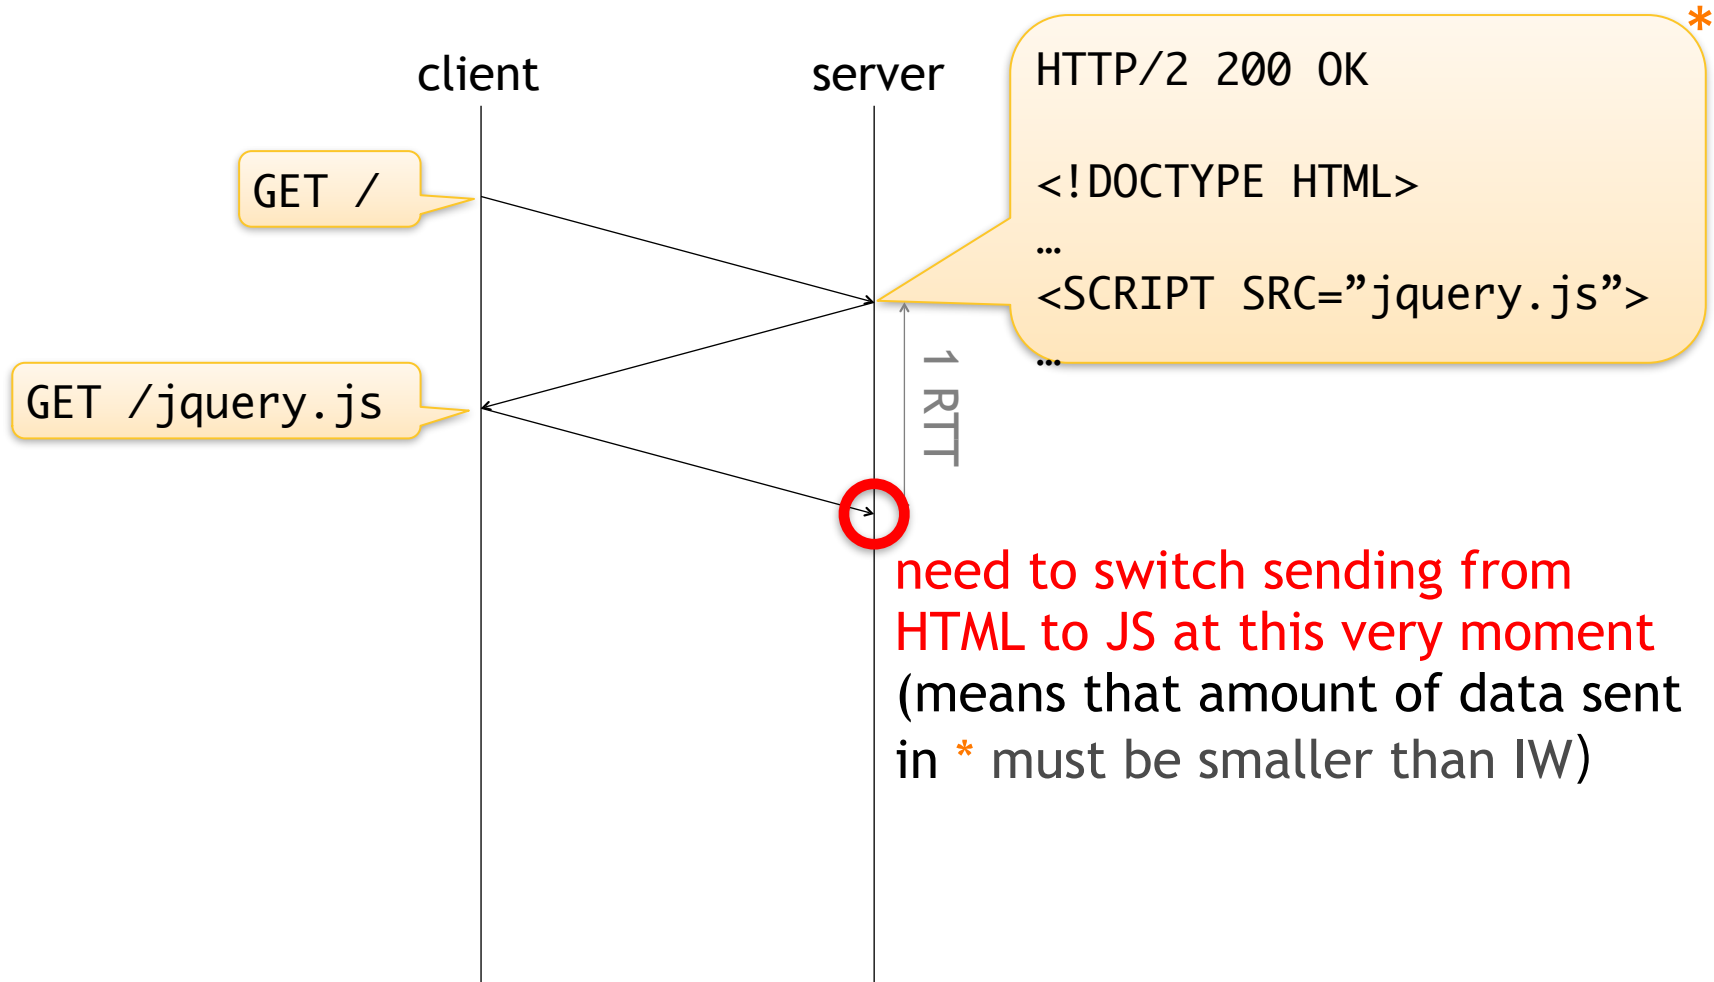
Programming TCP for responsiveness



DeNA Co., Ltd.


Kazuho Oku

Typical sequence of HTTP/2




HOB caused by buffers

- TCP send buffer:
 - to reduce ping-pong bet. kernel and application
- BIO buffer:
 - for data that couldn't be stored in TCP send buffer



Reduce poll threshold

- `setsockopt(TCP_NOTSENT_LOWAT)`
 - in linux, the minimum is $CWND + 1$ octet
 - becomes unstable when set to $CWND + 0$




How can we fill up the CWND?

- idea: do smaller writes until epoll tells you its full
- the issue with the idea:
 - CPU intensive
 - data overflowed from CWND might get sent as a small packet (wasting packets during slow start!)
 - overhead of TLS header & HTTP frame becomes bigger

Solution: read TCP states

```
// calc size of data to send by calling getsockopt(TCP_INFO)
if (poll_for_write(fd) == SOCKET_IS_READY) {
    capacity = CWND - unacked + TWO_MSS - TLS_overhead;
    SSL_write(prepare_http2_frames(capacity));
}
```



Negative impact of additional delay

- increased delay bet. ACK recv. → data send, since:
 - traditional approach: completes within kernel
 - this approach: application needs to be notified to generate new data
- outcome:
 - increase of CWND becomes slower
 - leads to slower peak speed?
 - depends on how CWND at peak is calculated
 - does kernel use TCP timestamp for the matter?

Countermeasures

- optimize for responsiveness only when necessary
 - i.e. when RTT is big *and* CWND is small
 - impact of optimization is $RTT * unsent_bytes / CWND$
- disable optimization if additional delay is significant
 - when epoll returns immediately, estimated additional delay is equal to the time spent by the loop

Configuration Directives

- http2-latopt-min-rtt
 - minimum TCP RTT to enable the optimization
 - default: UINT_MAX (disabled)
- http2-latopt-max-cwnd
 - maximum CWND to enable (in octets)
 - default: 65535
- http2-max-additional-delay
 - max. additional delay (as the ratio to TCP RTT)
 - latopt disabled if the delay is greater
 - default: 0.1

Pseudo-code

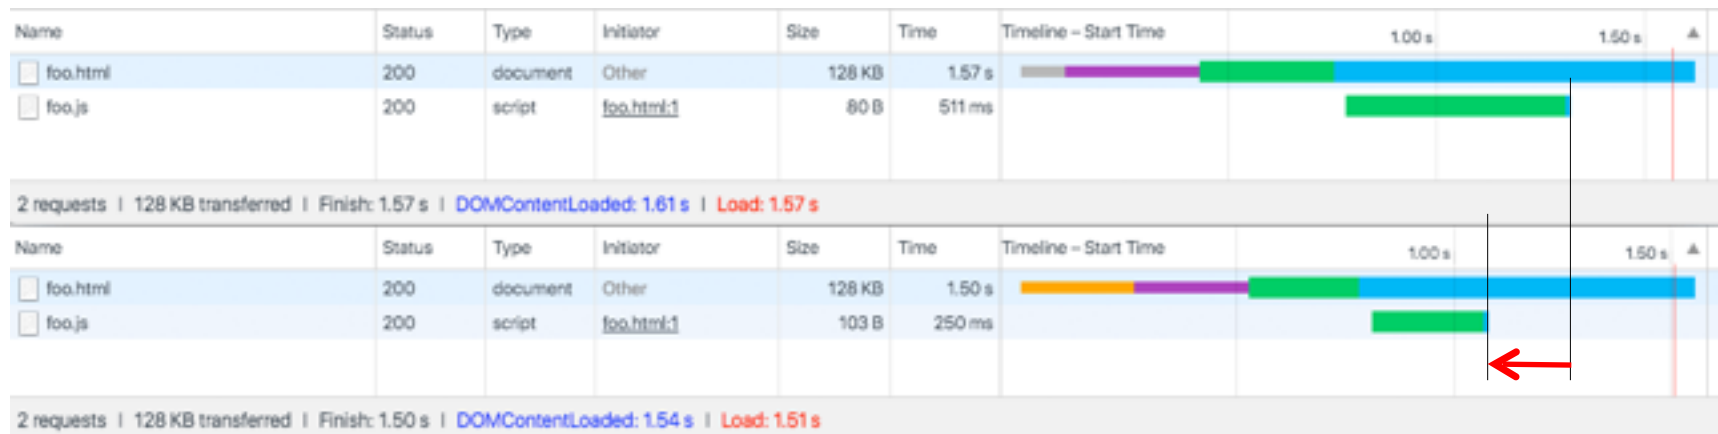
```
size_t get_suggested_write_size() {
    getsockopt(fd, IPPROTO_TCP, TCP_INFO, &tcp_info, sizeof(tcp_info));
    if (tcp_info.tcpi_rtt < min_rtt || tcp_info.tcpi_snd_cwnd > max_cwnd)
        return UNKNOWN;

    switch (SSL_get_current_cipher(ssl)->id) {
    case TLS1_CK_RSA_WITH_AES_128_GCM_SHA256:
    case ...:
        tls_overhead = 5 + 8 + 16;
        break;
    default:
        return UNKNOWN;
    }

    packets_sendable = tcp_info.tcpi_snd_cwnd > tcp_info.tcpi_unacked ?
        tcp_info.tcpi_snd_cwnd - tcp_info.tcpi_unacked : 0;
    return (packets_sendable + 2) * (tcp_info.tcpi_snd_mss - tls_overhead);
}
```


Benchmark (1)

- conditions:
 - server in Ireland, client in Tokyo (RTT 250ms)
 - load tiny js at the top of a large HTML
- result: delay decreased from 511ms to 250ms
 - i.e. JS fetch latency was 2RTT, became 1 RTT
 - similar results in other environments



Benchmark (2)

- using same data as previous
- server: Sakura VPS (Ishikari DC)



Conclusion

- near-optimal result can be achieved
 - by adjusting poll threshold and reading TCP states
 - 1-packet overhead due to restriction in Linux kernel
- 1-RTT improvement in H2O
 - estimated 1-RTT improvement per the depth of the load graph

Under the hood

TCP_NOTSENT_LOWAT

- supported by Linux, OS X
- on Linux:
 - sysctl:
 - set to -1: use kernel default
 - set to 0: sshd hangs
 - set to positive int: override kernel default
 - setsockopt:
 - set to 0: use default (sysctl or kernel)
 - set to int: override default

Unit of CWND

- Linux: # of packets
 - if INITCWND is 10, you can send at most 10 packets at once, regardless of their size
- BSD (incl. OS X): octets
 - you can send $CWND * MSS$ octets, regardless of the number of packets
 - if $CWND=10$ and $MSS=1460$, it is possible to send 14,600 packets containing 1-octet payload

Determining amount of data that can be sent immediately

- calculate either of:
 - $CWND - inflight$
 - $\min(CWND - (inflight + unsent), 0)$
- units used in the calculation must be the same
 - NetBSD: fail

OS	MSS	CWND	inflight	send buffer (inflight + unsent)
Linux	tcpi_snd_mss	tcpi_snd_cwnd*	tcpi_snd_unacked*	ioctl(SIOCOUTQ)
OS X**	tcpi_maxseg	tcpi_snd_cwnd	-	tcpi_snd_sbbytes
FreeBSD	tcpi_snd_mss	tcpi_snd_cwnd	-	ioctl(FIONWRITE)
*: units of values marked are packets, unmarked are octets				
**NetBSD: sometimes the values of tcpi_ are returned as zeros	tcpi_snd_mss	tcpi_snd_cwnd*		ioctl(FIONWRITE)